

```
In [23]: # Kemp Carswell 801017179
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [24]: df = pd.read_csv('C:/Users/kemp/Downloads/Housing.csv')
df.head() # To get first n rows from the dataset default value of n is 5
M=len(df)
```

```
In [25]: housing = pd.DataFrame(pd.read_csv('C:/Users/kemp/Downloads/Housing.csv'))
housing.head()
```

```
Out[25]:
```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterheating
0	13300000	7420	4	2	3	yes	no	no	no
1	12250000	8960	4	4	4	yes	no	no	no
2	12250000	9960	3	2	2	yes	no	yes	no
3	12215000	7500	4	2	2	yes	no	yes	no
4	11410000	7420	4	1	2	yes	yes	yes	no

```
In [26]: # You can see that your dataset has many columns with values as 'Yes' or 'No'.
# But in order to fit a regression line, we would need numerical values and not string.
# List of variables to map
varlist = ['mainroad', 'guestroom', 'basement', 'hotwaterheating', 'airconditioning', '']
# Defining the map function
def binary_map(x):
    return x.map({'yes': 1, "no": 0})
# Applying the function to the housing list
housing[varlist] = housing[varlist].apply(binary_map)
# Check the housing dataframe now
housing.head()
```

```
Out[26]:
```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterheating
0	13300000	7420	4	2	3	1	0	0	0
1	12250000	8960	4	4	4	1	0	0	0
2	12250000	9960	3	2	2	1	0	1	0
3	12215000	7500	4	2	2	1	0	1	0
4	11410000	7420	4	1	2	1	1	1	0

```
In [27]: #Splitting the Data into Training and Testing Sets
from sklearn.model_selection import train_test_split
# We specify this so that the train and test data set always have the same rows, respec
```

```
np.random.seed(0)
df_train, df_test = train_test_split(housing, train_size = 0.7, test_size = 0.3, random
```

```
In [28]: num_vars = ['area', 'bedrooms', 'bathrooms', 'stories', 'parking', 'price']
df_Newtrain = df_train[num_vars]
df_Newtest = df_test[num_vars]
df_Normalization = df_Newtrain
df_Standardization = df_Newtrain
df_Newtrain.head()
```

```
Out[28]:
```

	area	bedrooms	bathrooms	stories	parking	price
<b>454</b>	4500	3	1	2	0	3143000
<b>392</b>	3990	3	1	2	0	3500000
<b>231</b>	4320	3	1	1	0	4690000
<b>271</b>	1905	5	1	2	0	4340000
<b>250</b>	3510	3	1	3	0	4515000

```
In [29]: import warnings
warnings.filterwarnings('ignore')
from sklearn.preprocessing import MinMaxScaler, StandardScaler
# define standard scaler
#scaler = StandardScaler()
scaler = MinMaxScaler()
df_Normalization[num_vars] = scaler.fit_transform(df_Normalization[num_vars])
df_Normalization.head(20)
```

```
Out[29]:
```

	area	bedrooms	bathrooms	stories	parking	price
<b>454</b>	0.193548	0.50	0.0	0.333333	0.000000	0.120606
<b>392</b>	0.156495	0.50	0.0	0.333333	0.000000	0.151515
<b>231</b>	0.180471	0.50	0.0	0.000000	0.000000	0.254545
<b>271</b>	0.005013	1.00	0.0	0.333333	0.000000	0.224242
<b>250</b>	0.121622	0.50	0.0	0.666667	0.000000	0.239394
<b>541</b>	0.040976	0.50	0.0	0.000000	0.000000	0.001485
<b>461</b>	0.226969	0.25	0.0	0.000000	0.000000	0.115152
<b>124</b>	0.340671	0.50	0.5	1.000000	0.333333	0.363636
<b>154</b>	0.131793	0.50	0.5	0.333333	0.666667	0.327273
<b>451</b>	0.357018	0.25	0.0	0.000000	0.000000	0.121212
<b>59</b>	0.302528	0.50	0.5	1.000000	0.333333	0.472727
<b>493</b>	0.154316	0.50	0.0	0.000000	0.000000	0.090909
<b>465</b>	0.142691	0.25	0.0	0.000000	0.000000	0.112121
<b>490</b>	0.182650	0.50	0.0	0.333333	0.333333	0.093939

	area	bedrooms	bathrooms	stories	parking	price
<b>540</b>	0.084568	0.25	0.0	0.000000	0.666667	0.006061
<b>406</b>	0.253124	0.25	0.0	0.000000	0.333333	0.148485
<b>289</b>	0.291630	0.25	0.0	0.000000	0.666667	0.212121
<b>190</b>	0.418774	0.75	0.0	0.333333	0.666667	0.284848
<b>55</b>	0.302528	0.50	0.0	0.333333	0.333333	0.484848
<b>171</b>	0.612685	0.50	0.0	0.000000	0.333333	0.303030

In [30]:

```
import warnings
warnings.filterwarnings('ignore')
from sklearn.preprocessing import MinMaxScaler, StandardScaler

scaler = StandardScaler()
df_Standardization[num_vars] = scaler.fit_transform(df_Standardization[num_vars])
df_Standardization.head(20)
```

Out[30]:

	area	bedrooms	bathrooms	stories	parking	price
<b>454</b>	-0.286366	0.073764	-0.581230	0.207401	-0.822960	-0.868394
<b>392</b>	-0.544762	0.073764	-0.581230	0.207401	-0.822960	-0.677628
<b>231</b>	-0.377564	0.073764	-0.581230	-0.937813	-0.822960	-0.041744
<b>271</b>	-1.601145	2.884176	-0.581230	0.207401	-0.822960	-0.228768
<b>250</b>	-0.787958	0.073764	-0.581230	1.352614	-0.822960	-0.135256
<b>541</b>	-1.350349	0.073764	-0.581230	-0.937813	-0.822960	-1.603589
<b>461</b>	-0.053303	-1.331442	-0.581230	-0.937813	-0.822960	-0.902058
<b>124</b>	0.739618	0.073764	1.488383	2.497828	0.321375	0.631546
<b>154</b>	-0.717026	0.073764	1.488383	0.207401	1.465710	0.407116
<b>451</b>	0.853616	-1.331442	-0.581230	-0.937813	-0.822960	-0.864653
<b>59</b>	0.473622	0.073764	1.488383	2.497828	0.321375	1.304836
<b>493</b>	-0.559962	0.073764	-0.581230	-0.937813	-0.822960	-1.051678
<b>465</b>	-0.641027	-1.331442	-0.581230	-0.937813	-0.822960	-0.920761
<b>490</b>	-0.362365	0.073764	-0.581230	0.207401	0.321375	-1.032976
<b>540</b>	-1.046354	-1.331442	-0.581230	-0.937813	1.465710	-1.575348
<b>406</b>	0.129094	-1.331442	-0.581230	-0.937813	0.321375	-0.696331
<b>289</b>	0.397623	-1.331442	-0.581230	-0.937813	1.465710	-0.303578
<b>190</b>	1.284276	1.478970	-0.581230	0.207401	1.465710	0.145281
<b>55</b>	0.473622	0.073764	-0.581230	0.207401	0.321375	1.379646
<b>171</b>	2.636548	0.073764	-0.581230	-0.937813	0.321375	0.257496

```
In [31]: X_Training_N = df_Normalization.values[:,[0,1,2,3,4]]
         y_Training_N = df_Normalization.values[:,5]

         X_Test = df_Newtest.values[:,[0,1,2,3,4]]
         y_Test = df_Newtest.values[:,5]

         X_Training_S = df_Standardization.values[:,[0,1,2,3,4]]
         y_Training_S = df_Standardization.values[:,5]
```

```
In [21]: mean = np.ones(X_Training_N.shape[1])
         std = np.ones(X_Training_N.shape[1])
         for i in range(0, X_Training_N.shape[1]):
             mean[i] = np.mean(X_Training_N.transpose()[i])
             std[i] = np.std(X_Training_N.transpose()[i])
             for j in range(0, X_Training_N.shape[0]):
                 X_Training_N[j][i] = (X_Training_N[j][i] - mean[i])/std[i]
```

```
In [22]: mean = np.ones(X_Training_S.shape[1])
         std = np.ones(X_Training_S.shape[1])
         for i in range(0, X_Training_S.shape[1]):
             mean[i] = np.mean(X_Training_S.transpose()[i])
             std[i] = np.std(X_Training_S.transpose()[i])
             for j in range(0, X_Training_S.shape[0]):
                 X_Training_S[j][i] = (X_Training_S[j][i] - mean[i])/std[i]
```

```
In [32]: mean = np.ones(X_Test.shape[1])
         std = np.ones(X_Test.shape[1])
         for i in range(0, X_Test.shape[1]):
             mean[i] = np.mean(X_Test.transpose()[i])
             std[i] = np.std(X_Test.transpose()[i])
             for j in range(0, X_Test.shape[0]):
                 X_Test[j][i] = (X_Test[j][i] - mean[i])/std[i]
```

```
In [33]: def compute_cost(X, n, theta):
         h = np.ones((X.shape[0],1))
         theta = theta.reshape(1,n+1)
         for i in range(0,X.shape[0]):
             h[i] = float(np.matmul(theta, X[i]))
         h = h.reshape(X.shape[0])
         return h
```

```
In [34]: def gradient_descent(X, y, theta, alpha, iterations, n, h):
         cost = np.ones(iterations)
         for i in range(0,iterations):
             theta[0] = theta[0] - (alpha/X.shape[0]) * sum(h - y)
             for j in range(1,n+1):
                 theta[j] = theta[j] - (alpha/X.shape[0]) * sum((h-y) * X.transpose()[j])
             h = compute_cost(X, n, theta)
             cost[i] = (1/X.shape[0]) * 0.5 * sum(np.square(h - y))
             theta = theta.reshape(1,n+1)
         return theta, cost
```

```
In [35]: def linear_regression(X, y, alpha, iterations):
```

```

n = X.shape[1]
one_column = np.ones((X.shape[0],1))
X = np.concatenate((one_column, X), axis = 1)
theta = np.zeros(n+1)
h = compute_cost(X, n, theta)
theta, cost = gradient_descent(X, y, theta, alpha, iterations, n, h)
return theta, cost

```

```

In [36]: iterations = 500;
alpha = 0.01;
alpha2 = 0.01

```

```

In [37]: theta_Training, cost_Training = linear_regression(X_Training_N, y_Training_N, alpha, it
print('Final value of theta with normalization =', theta_Training)
cost_Training = list(cost_Training)
n_ierations_Training = [x for x in range(1,(iterations + 1))]

```

```

Final value of theta with normalization = [[8.12123772e-17  3.79061776e-01  1.10230128e-01
2.94608603e-01
2.32422016e-01  1.53858866e-01]]

```

```

In [38]: theta_Training2, cost_Training2 = linear_regression(X_Training_S, y_Training_S, alpha,
print('Final value of theta with standardization =', theta_Training2)
cost_Training2 = list(cost_Training2)
n_ierations_Training2 = [x for x in range(1,(iterations + 1))]

```

```

Final value of theta with standardization = [[8.12123772e-17  3.79061776e-01  1.10230128e-
01  2.94608603e-01
2.32422016e-01  1.53858866e-01]]

```

```

In [39]: theta_Test, cost_Test = linear_regression(X_Test, y_Test, alpha, iterations)
print('Final value of theta of the test set =', theta_Test)
cost_Test = list(cost_Test)
n_ierations_Test = [x for x in range(1,(iterations + 1))]

```

```

Final value of theta of the test set = [[3896885.81334708  798864.59108174  151510.77459
081  1093108.39710527
870883.11233557  848681.31817011]]

```

```


In [40]: plt.plot(n_ierations_Training, cost_Training, label='Traing set normalization')
plt.legend()
plt.rcParams["figure.figsize"]=(10,6)
plt.grid()
plt.xlabel('No. of iterations')
plt.ylabel('Cost')
plt.title('Convergence of gradient descent')

```

```


Out[40]: Text(0.5, 1.0, 'Convergence of gradient descent')

```




```
In [41]: plt.plot(n_ierations_Training2, cost_Training2, label='Traing set standardization')
plt.legend()
plt.rcParams["figure.figsize"]=(10,6)
plt.grid()
plt.xlabel('No. of iterations')
plt.ylabel('Cost')
plt.title('Convergence of gradient descent')
```

Out[41]: Text(0.5, 1.0, 'Convergence of gradient descent')



```
In [42]: plt.plot(n_ierations_Test, cost_Test, label='Test set')
plt.legend()
plt.rcParams["figure.figsize"]=(10,6)
plt.grid()
plt.xlabel('No. of iterations')
plt.ylabel('Cost')
plt.title('Convergence of gradient descent')
```

Out[42]: Text(0.5, 1.0, 'Convergence of gradient descent')



In [ ]: